Dynamic Trend & Event Detector

Emerging Topic Detection & News Correlation

Deep Learning & Advanced Machine Learning — 6th Semester

Pintu Singh (230105) Fathal (230043)

March 2026

Contents

1 Abstract

This project presents a Dynamic Trend and Event Detector that combines Advanced Machine Learning (Latent Dirichlet Allocation) with Deep Learning (BERTopic using SBERT embeddings) to detect and track emerging topics from news data streams.

The core challenge is distinguishing between fleeting viral trends and significant real-world news events. We compare three approaches — Word Frequency baseline, LDA (Advanced ML), and BERTopic (Deep Learning). BERTopic achieves the best coherence score of 0.4781, improving 33.8% over the LDA baseline of 0.3573.

Keywords: Topic Modeling, LDA, BERTopic, SBERT, UMAP, HDBSCAN, NLP, Unsupervised Learning

2 Literature Review

2.1 Problem Gap

Organizations monitor massive news streams daily. Existing approaches treat documents as static collections rather than evolving streams. They cannot distinguish between:

- **Real breaking events:** COVID topic grew from 50 articles in January 2020 to 5,000 articles in March 2020
- **Viral noise:** A meme viral for one day then disappears completely

This project addresses this gap by combining probabilistic topic modeling with neural semantic embeddings.

2.2 Simple Frequency-Based Extraction

The simplest baseline counts word occurrences:

$$\text{Frequency}(w) = \sum_{d \in D} \text{count}(w, d) \tag{1}$$

Example from our dataset:

Word	Frequency
trump	5,420
health	3,210
police	2,890
election	2,100
climate	1,980

Limitation: trump, election, vote counted separately — no grouping into one Politics topic.

2.3 Latent Dirichlet Allocation (LDA)

Blei, Ng, and Jordan [1] introduced LDA as a generative probabilistic model:

$$p(\theta, z, w|\alpha, \beta) = p(\theta|\alpha) \prod_{n=1}^{N} p(z_n|\theta) \cdot p(w_n|z_n, \beta) \tag{2}$$

Where:

- θ = topic distribution per document
- z = topic assignment per word
- w = observed words
- α, β = Dirichlet hyperparameters

Example: One article represented as: 70% Politics + 20% Economy + 10% Health.

Critical limitation: Bag-of-words — covid and virus treated as completely different words despite same meaning.

2.4 BERTopic

Grootendorst [2] proposed BERTopic using a three-stage neural pipeline:

1. **SBERT** — semantic embeddings that capture word meaning
2. **UMAP** — dimensionality reduction from 384 to 5 dimensions
3. **HDBSCAN** — automatic density based clustering

Key advantage: covid and virus get similar vector representations — semantic meaning captured unlike LDA.

2.5 Gap This Project Fills

Method	Groups Words	Semantic	Auto Topics
Word Frequency	No	No	No
LDA (Phase 1)	Yes	No	No
BERTopic (Phase 2)	Yes	Yes	Yes

Table 1: Comparison of approaches used

BERTopic is the only approach that captures semantic meaning and detects topics automatically.

3 Dataset & Exploratory Data Analysis

3.1 Dataset Description

Property	Value
Dataset Name	News Category Dataset
Source	Kaggle / HuffPost
Original Size	209,527 articles
Sampled Size	10,000 articles
Date Range	January 2012 — September 2022
Categories	41 unique
Format	JSON (lines=True)

Table 2: Dataset properties

3.2 EDA Finding 1: Class Imbalance

Category	Count	Percentage
POLITICS	3,298	33.0%
WELLNESS	1,794	17.9%
ENTERTAINMENT	1,736	17.4%
TRAVEL	961	9.6%
Others	2,211	22.1%

Action taken: Unsupervised topic modeling chosen — model discovers topics without label bias.

3.3 EDA Finding 2: Non-Stationarity

Year	Articles
2012	412
2015	987
2018	1,891
2020	2,134
2022	876

2020 had highest volume due to COVID-19 and US Elections — non-stationary data. *Action taken:* Time-based analysis applied.

3.4 EDA Finding 3: Short Text

Before combining:

Headline only: "Trump Wins Presidential Election" (9 words average)

After combining headline and description:

> “Trump Wins Presidential Election. Donald Trump secured enough electoral votes to win the presidency, defeating Hillary Clinton in a historic upset.” (22 words average)

Action taken: Concatenated headline and short_description — average word count increased from 9 to 22 words.

3.5 Top Words After Cleaning

Word	**Frequency**
trump	5,420
people	4,210
health	3,210
state	3,100
election	2,100

4 Feature Engineering

4.1 Text Preprocessing

Six cleaning steps applied to every article:

1. Lowercase conversion
2. URL removal using regex
3. Non-alphabetic character removal
4. Stop word removal (NLTK English stopwords)
5. Lemmatization using WordNetLemmatizer
6. Removal of words under 3 characters

Before Cleaning	**After Cleaning**
"The President is running!"	"president run"
"COVID-19 cases rising"	"covid case rise"
"https://news.com/article"	"article"

4.2 Phase 1 Features: Word Frequency and TF-IDF

Word Frequency using Python Counter:

```
Counter({'trump': 5420, 'health': 3210,
         'election': 2100, 'climate': 1980})
```

TF-IDF for LDA using CountVectorizer: `max_features=10000`, `min_df=5`, `max_df=0.95`

$$\text{TF-IDF}(t, d) = \text{TF}(t, d) \times \log\left(\frac{N}{\text{DF}(t)}\right) \tag{3}$$

Example: The word "the" appears in all documents so its weight $= 0$ automatically. The word "election" appears in fewer documents so it gets higher weight.

4.3 Phase 2 Features: SBERT Embeddings

For BERTopic, `all-MiniLM-L6-v2` produces 384-dimensional dense vectors:

Article	**Vector**
"trump wins election"	[0.21, 0.84, 0.12, ...] (384 dims)
"biden wins election"	[0.23, 0.81, 0.14, ...] (384 dims)
"covid cases rising"	[0.91, 0.02, 0.76, ...] (384 dims)

First two articles have similar vectors — semantic similarity successfully captured!

5 Theoretical Rigor

5.1 Why Unsupervised Learning

Dataset has 41 category labels but goal is to discover fine-grained sub-topics:

- Category label = POLITICS
- Sub-topics inside POLITICS:
  - Trump election campaign
  - Senate voting bills
  - Foreign policy debates
  - Immigration discussions

Supervised models (SVM, Classification) predict POLITICS label only — cannot discover sub-topics. Unsupervised models (LDA, BERTopic) discover all sub-topics automatically without labels.

5.2 Why LDA Over K-Means

Aspect	**K-Means**	**LDA**
K required	Yes (must specify)	No
Cluster shape	Spherical only	Any shape
Outliers	Cannot handle	Handles well

We do not know how many topics exist in 10,000 articles — K-Means would require guessing this number upfront.

5.3 Why BERTopic Over LDA

Aspect	**LDA**	**BERTopic**
Word meaning	Ignored	Captured (SBERT)
Topic count	Fixed (20)	Automatic (28)
Clustering	Probabilistic	Density-based
Coherence	0.3573	0.4781

5.4 Bias-Variance Analysis

- **Word Frequency (Baseline):** Extremely high bias. Example: trump and election counted separately — no insight.
- **LDA (Phase 1):** Lower bias — groups related words. Example: trump, election, vote grouped into Topic 0 Politics. Controlled variance via random_state=42.
- **BERTopic (Phase 2):** Lowest bias — semantic meaning captured. Example: covid and virus automatically placed in same topic.

5.5 LDA Assumptions and Violations

1. **IID assumption violated:** Articles from same news cycle are correlated. Example: 100 COVID articles in March 2020 are all related — not independent.

2. **Bag-of-words violated:** covid and coronavirus have same meaning but LDA treats them as different words.

These violations explain moderate score of 0.3573 and motivated BERTopic as Phase 2 improvement.

6 Model Application

6.1 Phase 1 — Baseline: Word Frequency

- Tool: Python Counter()
- Top word: trump (5,420 occurrences)
- No coherence score applicable
- Cannot group related words into topics

Limitation: trump, election, vote counted separately — no way to know they form one Politics topic.

6.2 Phase 1 — Model A: LDA (Advanced ML)

- Vectorizer: CountVectorizer 10,000 features
- Topics: 20 fixed via n_components
- Iterations: 20, online learning
- Random state: 42 (reproducible)
- Coherence Score (c_v): **0.3573**

Topics discovered by LDA:

Topic	Top Words	Label
0	trump, election, vote, president	Politics
1	health, hospital, disease, medical	Health
2	climate, carbon, energy, policy	Environment
3	game, team, player, season	Sports
4	film, movie, actor, award	Entertainment

LDA Limitation identified: covid and virus appeared in different topics despite having the same meaning. This directly motivated Phase 2 — BERTopic.

6.3 Phase 2 — Model B: BERTopic (Deep Learning)

- Embedding: SBERT all-MiniLM-L6-v2
- Reduction: UMAP 384 $\rightarrow$ 5 dimensions
- Clustering: HDBSCAN min_cluster_size=15
- Topics: 28 auto-detected (no k needed!)
- Coherence Score (c_v): **0.4781**
- Improvement over LDA: +**33.8%**

BERTopic improvement over LDA: covid and virus now appear in the same topic because SBERT understands their similar meaning.

6.4 Phase 1 vs Phase 2 Comparison

Model	Phase	Type	Score	Semantic
Word Frequency	1	Baseline	N/A	No
LDA	1	Advanced ML	0.3573	No
BERTopic	2	Deep Learning	0.4781	Yes

Table 3: Phase 1 vs Phase 2 — BERTopic wins

Conclusion: BERTopic (Phase 2) outperforms LDA (Phase 1) by 33.8% because it understands semantic meaning through SBERT transformer embeddings.

7 Results & Discussion

7.1 Phase 1 Results — LDA

LDA successfully discovered 20 meaningful topics. Coherence score of 0.3573 confirms topics are meaningful — top words co-occur in real articles.

Limitation found: covid and virus appeared in different topics because LDA uses bag-of-words and ignores meaning.

7.2 Phase 2 Results — BERTopic

BERTopic automatically discovered 28 topics without specifying k in advance. Coherence score improved to 0.4781.

Model	**Score**	**Improvement**
Word Frequency	N/A	Baseline
LDA	0.3573	—
BERTopic	0.4781	+33.8%

7.3 Why BERTopic Won

1. **Semantic understanding:** SBERT embeddings capture word meaning. covid and virus get similar vectors — placed in same topic automatically.
2. **Automatic topic count:** LDA required manually setting 20 topics. BERTopic found 28 topics automatically using HDBSCAN density clustering.
3. **Better representation:** 384-dimensional SBERT vectors capture much richer information than TF-IDF sparse vectors used by LDA.

7.4 Future Work — Phase 3

These results motivate the following extensions:

- **Temporal Tracking:** Track how topics evolve over time. Detect when a topic spikes — breaking event!
- **Semantic Velocity:** Measure how fast a topic grows. High velocity = breaking news event.
- **GDELT Verification:** Validate detected topics against real-world news database externally.

8 Conclusion

This project successfully demonstrated that:

1. **Word Frequency (Phase 1 Baseline)** counts words but cannot group related words. trump and election counted separately — no meaningful topics formed.

2. **LDA (Phase 1 Advanced ML)** discovers 20 meaningful topics automatically without any human labeling. Coherence score: 0.3573. Limitation: bag-of-words ignores meaning.

3. **BERTopic (Phase 2 Deep Learning)** improves coherence by 33.8% over LDA by understanding semantic meaning through SBERT transformer embeddings. Automatically found 28 topics — no k needed.

4. **Deep Learning beats Advanced ML** — proven through coherence score comparison. SBERT embeddings are clearly superior to TF-IDF for topic modeling tasks.

References

[1] Blei, D.M., Ng, A.Y., Jordan, M.I. (2003). *Latent Dirichlet Allocation*. Journal of Machine Learning Research, 3, 993–1022.

[2] Grootendorst, M. (2022). *BERTopic: Neural topic modeling with a class-based TF-IDF procedure*. arXiv preprint arXiv:2203.05794.

[3] Reimers, N., Gurevych, I. (2019). *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks*. EMNLP 2019.

[4] McInnes, L., Healy, J. (2017). *HDBSCAN: Hierarchical density based clustering*. Journal of Open Source Software, 2(11), 205.

[5] Misra, R. (2022). *News Category Dataset*. Kaggle.

[6] Pedregosa et al. (2011). *Scikit-learn: Machine Learning in Python*. JMLR, 12, 2825–2830.

[7] Bird, S., Klein, E., Loper, E. (2009). *Natural Language Processing with Python*. O'Reilly Media.