Project : Dynamic Trend & Event Detector

Emerging Topic Detection & News Correlation

Deep Learning & Advanced Machine Learning — 6th Semester

Authors: Pintu Singh (230105) Fathal (230043)

March 22, 2026

Contents

1 Abstract **2**

2 Literature Review **3**
2.1 Problem Gap . . . 3
2.2 Latent Dirichlet Allocation (LDA) . . . 3
2.3 BERTopic . . . 3
2.4 Gap This Project Fills . . . 3

3 Dataset & Exploratory Data Analysis **4**
3.1 Dataset Description . . . 4
3.2 EDA Findings . . . 4

4 Feature Engineering **5**
4.1 Baseline Features: TF-IDF . . . 5
4.2 Advanced Features: SBERT Embeddings . . . 5
4.3 Novel Feature: Semantic Velocity . . . 5

5 Theoretical Rigor **6**
5.1 Why LDA Fails for This Task . . . 6
5.2 Why HDBSCAN Over K-Means . . . 6
5.3 Bias-Variance Analysis . . . 6

6 Model Application **7**
6.1 Model A — LDA (Advanced ML Baseline) . . . 7
6.2 Model B — BERTopic (Deep Learning) . . . 7
6.3 Model C — Hybrid System . . . 7
6.4 Ablation Table . . . 7

7 Results & Validation **8**
7.1 GDELT External Verification . . . 8
7.2 Semantic Velocity Findings . . . 8

8 Conclusion **9**

1 Abstract

This project presents a Hybrid Dynamic Trend and Event Detector that combines Advanced Machine Learning (Latent Dirichlet Allocation) with Deep Learning (BERTopic using SBERT embeddings) to detect, track, and validate emerging topics from news data streams. The core challenge addressed is distinguishing between fleeting social media memes and significant real-world news events by tracking semantic velocity over time. Our system demonstrates that the Hybrid approach achieves superior topic coherence (0.4981) compared to the LDA baseline (0.3573), validated against the GDELT news database.

Keywords: Topic Modeling, BERTopic, LDA, SBERT, Temporal Trend Detection, GDELT, NLP

2 Literature Review

2.1 Problem Gap

Existing topic modeling approaches face a fundamental limitation: they treat documents as static collections rather than evolving streams. Traditional bag-of-words methods lose temporal dynamics and semantic relationships between words. This project addresses this specific gap by combining probabilistic topic modeling with neural embeddings to track how topics evolve over time.

2.2 Latent Dirichlet Allocation (LDA)

Blei, Ng, and Jordan [1] introduced LDA as a generative probabilistic model where each document is modeled as a mixture of topics, and each topic is a distribution over words. Mathematically:

$$p(\theta, z, w|\alpha, \beta) = p(\theta|\alpha) \prod_{n=1}^{N} p(z_n|\theta) \cdot p(w_n|z_n, \beta) \tag{1}$$

Critical limitation: LDA assumes word independence (bag-of-words) and ignores semantic similarity. "Car" and "automobile" are treated as completely different words, causing topic fragmentation.

2.3 BERTopic

Grootendorst [2] proposed BERTopic, which solves LDA's limitations through a three-stage pipeline:

1. **SBERT embeddings** — captures semantic meaning
2. **UMAP** — reduces 384 dimensions to 5
3. **HDBSCAN** — density-based clustering

2.4 Gap This Project Fills

Method	Semantic	Temporal	Auto Topics
LDA (Blei 2003)	No	No	No
BERTopic (2022)	Yes	No	Yes
Our Hybrid	Yes	Yes	Yes

Table 1: Comparison of topic modeling approaches

3 Dataset & Exploratory Data Analysis

3.1 Dataset Description

Property	Value
Dataset Name	News Category Dataset
Source	Kaggle / HuffPost
Original Size	209,527 articles
Sampled Size	10,000 articles
Date Range	January 2012 — September 2022
Categories	41 unique categories

Table 2: Dataset properties

3.2 EDA Findings

Three key findings emerged from exploratory data analysis:

1. **Class Imbalance:** POLITICS category dominates with 33% of articles. *Action taken:* Unsupervised topic modeling — no labels used.

2. **Non-Stationarity:** Article volume varies 3x across years (peak in 2017–2019). *Action taken:* Time-binned analysis using 20 temporal windows.

3. **Short Text:** Mean headline length is 9.2 words. *Action taken:* Concatenated headline and short_description for richer semantic content.

4 Feature Engineering

4.1 Baseline Features: TF-IDF

For the LDA baseline, we use CountVectorizer with: `max_features=10000`, `min_df=5`, `max_df=0.95`.

4.2 Advanced Features: SBERT Embeddings

For BERTopic, we use `all-MiniLM-L6-v2` Sentence-BERT which produces 384-dimensional dense vectors per article.

4.3 Novel Feature: Semantic Velocity

$$\text{Velocity}(t) = \frac{\sigma(\text{Frequency}_t)}{\mu(\text{Frequency}_t) + 1} \tag{2}$$

High velocity = rapidly changing topic = breaking news event. Low velocity = stable background topic.

5 Theoretical Rigor

5.1 Why LDA Fails for This Task

LDA violates two assumptions in our dataset:

- **IID assumption violated:** Articles from the same news cycle are correlated.
- **Bag-of-words:** Semantic similarity ignored.

5.2 Why HDBSCAN Over K-Means

K-Means requires k to be pre-specified and assumes spherical clusters. News topics violate both — HDBSCAN uses density-based clustering with no k required.

5.3 Bias-Variance Analysis

- **LDA:** High bias (strong assumptions), low variance
- **BERTopic:** Lower bias, controlled variance via `random_state=42`
- **Hybrid:** Balanced — best of both

6 Model Application

6.1 Model A — LDA (Advanced ML Baseline)

- Vectorizer: CountVectorizer, 10,000 features
- Topics: 20 (fixed)
- Training: 20 iterations, online learning
- Coherence Score (c_v): **0.3573**

6.2 Model B — BERTopic (Deep Learning)

- Embeddings: SBERT all-MiniLM-L6-v2
- Reduction: UMAP (384 → 5 dimensions)
- Clustering: HDBSCAN (min_cluster_size=15)
- Topics: Auto-detected
- Coherence Score (c_v): **0.4781**

6.3 Model C — Hybrid System

- Base: BERTopic
- Addition 1: Temporal tracking (20 time bins)
- Addition 2: Semantic Velocity calculation
- Addition 3: GDELT external verification
- Coherence Score (c_v): **0.4981**

6.4 Ablation Table

Model	Type	Coherence	Temporal	GDELT
A — LDA	Advanced ML	0.3573	No	No
B — BERTopic	Deep Learning	0.4781	No	No
C — Hybrid	Hybrid	0.4981	Yes	Yes

Table 3: Ablation study results — Model C outperforms both baselines

Conclusion: Hybrid system improved coherence by **39.4%** over LDA baseline.

7 Results & Validation

7.1 GDELT External Verification

Topic 0 (Politics: trump, election, president) was successfully verified against GDELT real-world news database with 5 matching articles. This confirms our system detects real events, not noise.

7.2 Semantic Velocity Findings

High-velocity topics corresponded to breaking news events (elections, policy announcements). Low-velocity topics corresponded to stable background issues (health, environment).

8 Conclusion

This project successfully demonstrated that:

1. LDA (Advanced ML) provides a solid baseline but fails to capture semantic relationships.
2. BERTopic (Deep Learning) improves coherence by 33.8% over LDA by using SBERT transformer embeddings.
3. The Hybrid system adds temporal tracking and external validation, achieving the highest coherence score.
4. Semantic Velocity successfully distinguishes breaking news events from stable background topics.

References

[1] Blei, D.M., Ng, A.Y., Jordan, M.I. (2003). *Latent Dirichlet Allocation*. Journal of Machine Learning Research, 3, 993–1022.

[2] Grootendorst, M. (2022). *BERTopic: Neural topic modeling with a class-based TF-IDF procedure*. arXiv preprint arXiv:2203.05794.

[3] Reimers, N., Gurevych, I. (2019). *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks*. EMNLP 2019.

[4] McInnes, L., Healy, J. (2017). *HDBSCAN: Hierarchical density based clustering*. Journal of Open Source Software, 2(11), 205.

[5] Misra, R. (2022). *News Category Dataset*. Kaggle.